

```
In [60]: import pandas as pd
import warnings
warnings.filterwarnings("ignore")
```

```
In [61]: data=pd.read_csv("/home/placement/Downloads/fiat500.csv")
```

```
In [62]: data.describe()
```

```
Out[62]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [63]: data1=data.loc[(data.previous_owners)==1]
```

In [64]: data1

Out[64]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1389 rows × 9 columns

In [65]: data2=data1.drop(['lat', 'lon', 'ID'],axis=1)

```
In [66]: data2
```

```
Out[66]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1389 rows × 6 columns

```
In [67]: data2=pd.get_dummies(data2)
```

```
In [68]: data2
```

```
Out[68]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1389 rows × 8 columns

```
In [69]: y=data2['price']  
x=data2.drop('price',axis=1)
```

In [70]:

y

Out[70]:

0	8900
1	8800
2	4200
3	6000
4	5700
	...
1533	5200
1534	4600
1535	7500
1536	5990
1537	7900

Name: price, Length: 1389, dtype: int64

In [71]:

```
from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(x_train,y_train)
```

Out[71]:

LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [72]:

LinearRegression()

Out[72]:

LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [73]:

```
#!/pip3 install scikit-learn
```

In [74]:

```
ypred=reg.predict(x_test)
```

In [75]: ypred

Out[75]: array([5481.93168764, 5127.11081209, 4798.43164854, 9659.36578585,
 9409.4127446 , 10351.98379749, 9802.72406141, 8334.75329195,
 5913.57169572, 10150.04762334, 5643.36202062, 7780.90416594,
 9721.15872463, 4456.3882388 , 6541.53947176, 9829.09275112,
 7574.52796156, 5909.39873877, 10416.87928247, 7409.77542821,
 8693.13864599, 8182.36608361, 9441.1300824 , 10383.66774161,
 9857.9433171 , 10388.58335816, 9818.87050889, 7023.92041959,
 9335.62476174, 10173.88293864, 5551.06753428, 9769.38528629,
 4609.76045054, 9962.4794893 , 9789.3539293 , 8904.50209071,
 3336.10690574, 10067.44590413, 8607.43409685, 7682.12076521,
 10206.23086655, 10451.29193617, 10428.25147613, 9711.27231338,
 9296.17132987, 7217.0720428 , 10459.74879956, 9083.60035288,
 10416.67497977, 8567.06083756, 10390.98325814, 7953.60968003,
 5590.45997234, 10404.33169149, 5658.96046682, 8904.50209071,
 9962.4794893 , 5204.32975664, 8381.41911545, 6642.92293048,
 6236.53789235, 4815.11945754, 10356.87473279, 7963.88315168,
 5015.51747675, 9896.61284815, 8728.78349613, 5415.22108385,
 9921.17107046, 7314.69366999, 10088.79553655, 8210.01253214,
 10343.75594017, 10399.71785545, 9720.01037852, 9579.33859859,
 10477.05110102, 8047.52420100, 10172.14440204, 8666.52225060])

In [76]: `from sklearn.metrics import r2_score`
`r2_score(y_test,ypred)`

Out[76]: 0.8601937431943691

In [77]: `from sklearn.metrics import mean_squared_error as ns`
`o=ns(y_test,ypred)`
`o`

Out[77]: 515432.90107231616

In [78]: `import math`
`math.sqrt(o)`

Out[78]: 717.9365578324565

In [79]: ypred

Out[79]: array([5481.93168764, 5127.11081209, 4798.43164854, 9659.36578585,
 9409.4127446 , 10351.98379749, 9802.72406141, 8334.75329195,
 5913.57169572, 10150.04762334, 5643.36202062, 7780.90416594,
 9721.15872463, 4456.3882388 , 6541.53947176, 9829.09275112,
 7574.52796156, 5909.39873877, 10416.87928247, 7409.77542821,
 8693.13864599, 8182.36608361, 9441.1300824 , 10383.66774161,
 9857.9433171 , 10388.58335816, 9818.87050889, 7023.92041959,
 9335.62476174, 10173.88293864, 5551.06753428, 9769.38528629,
 4609.76045054, 9962.4794893 , 9789.3539293 , 8904.50209071,
 3336.10690574, 10067.44590413, 8607.43409685, 7682.12076521,
 10206.23086655, 10451.29193617, 10428.25147613, 9711.27231338,
 9296.17132987, 7217.0720428 , 10459.74879956, 9083.60035288,
 10416.67497977, 8567.06083756, 10390.98325814, 7953.60968003,
 5590.45997234, 10404.33169149, 5658.96046682, 8904.50209071,
 9962.4794893 , 5204.32975664, 8381.41911545, 6642.92293048,
 6236.53789235, 4815.11945754, 10356.87473279, 7963.88315168,
 5015.51747675, 9896.61284815, 8728.78349613, 5415.22108385,
 9921.17107046, 7314.69366999, 10088.79553655, 8210.01253214,
 10343.75594017, 10399.71785545, 9720.01037852, 9579.33859859,
 10477.05110102, 8047.52420100, 10172.14440204, 8666.53225060])

```
In [80]: Results=pd.DataFrame(columns=['price','predicted'])
Results['price']=y_test
Results['predicted']=ypred
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(15)
```

```
Out[80]:
```

	index	price	predicted	ID
0	625	5400	5481.931688	0
1	187	5399	5127.110812	1
2	279	4900	4798.431649	2
3	734	10500	9659.365786	3
4	315	9300	9409.412745	4
5	652	10850	10351.983797	5
6	1472	9500	9802.724061	6
7	619	7999	8334.753292	7
8	992	6300	5913.571696	8
9	1154	10000	10150.047623	9
10	757	6000	5643.362021	10
11	1299	8500	7780.904166	11
12	400	8580	9721.158725	12
13	314	4600	4456.388239	13
14	72	7400	6541.539472	14

```
In [81]: Results['price_diff']=Results.apply(lambda row: row.price - row.predicted,axis=1)
```


In [82]: Results

Out[82]:

	index	price	predicted	ID	price_diff
0	625	5400	5481.931688	0	-81.931688
1	187	5399	5127.110812	1	271.889188
2	279	4900	4798.431649	2	101.568351
3	734	10500	9659.365786	3	840.634214
4	315	9300	9409.412745	4	-109.412745
...
454	115	10650	10397.402425	454	252.597575
455	370	9900	10231.829592	455	-331.829592
456	1179	5900	6764.023619	456	-864.023619
457	93	10050	10378.419299	457	-328.419299
458	147	9900	10070.703624	458	-170.703624

459 rows × 5 columns

```
In [83]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
#ridge regression
alpha = [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 30]

ridge = Ridge()

parameters = {'alpha': alpha}

ridge_regressor = GridSearchCV(ridge, parameters)

ridge_regressor.fit(x_train, y_train)
```

```
Out[83]: GridSearchCV(estimator=Ridge(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20, 30]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [84]: ridge_regressor.best_params_
```

```
Out[84]: {'alpha': 20}
```

```
In [85]: ridge=Ridge(alpha=30)
ridge.fit(x_train,y_train)
y_pred_ridge=ridge.predict(x_test)
```

```
In [86]: from sklearn.metrics import mean_squared_error
Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
Ridge_Error
```

```
Out[86]: 515419.9621427437
```

```
In [87]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_ridge)
```

```
Out[87]: 0.8601972527555687
```

```
In [88]: Results=pd.DataFrame(columns=['Actual','predicted'])
Results['Actual']=y_test
Results['predicted']=y_pred_ridge
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(10)
```

```
Out[88]:
```

	index	Actual	predicted	ID
0	625	5400	5480.612378	0
1	187	5399	5126.772562	1
2	279	4900	4823.164641	2
3	734	10500	9679.384113	3
4	315	9300	9404.679979	4
5	652	10850	10346.266387	5
6	1472	9500	9822.477584	6
7	619	7999	8367.522197	7
8	992	6300	5912.518318	8
9	1154	10000	10144.696863	9

```
In [89]: from sklearn.linear_model import ElasticNet
from sklearn.model_selection import GridSearchCV

elastic = ElasticNet()

parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}

elastic_regressor = GridSearchCV(elastic, parameters)

elastic_regressor.fit(x_train, y_train)
```

```
Out[89]: GridSearchCV(estimator=ElasticNet(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [90]: elastic_regressor.best_params_
```

```
Out[90]: {'alpha': 0.01}
```

```
In [91]: elastic=ElasticNet(alpha=.01)
elastic.fit(x_train,y_train)
y_pred_elastic=elastic.predict(x_test)
```

```
In [92]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_elastic)
```

```
Out[92]: 0.8602162350730707
```

```
In [93]: from sklearn.metrics import mean_squared_error
elastic_Error=mean_squared_error(y_pred_elastic,y_test)
elastic_Error
```

```
Out[93]: 515349.9787871871
```

```
In [94]: Results=pd.DataFrame(columns=['Actual','predicted'])
Results['Actual']=y_test
Results['predicted']=y_pred_elastic
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(10)
```

```
Out[94]:
```

	index	Actual	predicted	ID
0	625	5400	5482.171479	0
1	187	5399	5127.531740	1
2	279	4900	4803.203231	2
3	734	10500	9662.825235	3
4	315	9300	9408.645424	4
5	652	10850	10350.952605	5
6	1472	9500	9806.127960	6
7	619	7999	8341.142824	7
8	992	6300	5913.786719	8
9	1154	10000	10149.093829	9

```
In [ ]:
```