

```
In [4]: import pandas as pd
```

```
In [5]: data=pd.read_csv("/home/placement/Desktop/naren/fiat500.csv")
```

```
#data['model']=data['model'].map({'lounge':1,'pop':2,'sport':3})
```

```
In [6]: data1=data.drop(['ID','lat','lon'],axis=1)
```

```
In [7]: data1
```

```
Out[7]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

Type *Markdown* and LaTeX: α^2

```
In [8]: #data1=pd.get_dummies(data1)
```

```
In [9]: #data1
```

```
In [10]: data1['model']=data1['model'].map({'lounge':1,'pop':2,'sport':3})
```

```
In [11]: data1
```

```
Out[11]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	1	51	882	25000	1	8900
1	2	51	1186	32500	1	8800
2	3	74	4658	142228	1	4200
3	1	51	2739	160000	1	6000
4	2	73	3074	106880	1	5700
...
1533	3	51	3712	115280	1	5200
1534	1	74	3835	112000	1	4600
1535	2	51	2223	60457	1	7500
1536	1	51	2557	80750	1	5990
1537	2	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [12]: data1=pd.get_dummies(data1)
```

```
In [13]: data1
```

```
Out[13]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	1	51	882	25000	1	8900
1	2	51	1186	32500	1	8800
2	3	74	4658	142228	1	4200
3	1	51	2739	160000	1	6000
4	2	73	3074	106880	1	5700
...
1533	3	51	3712	115280	1	5200
1534	1	74	3835	112000	1	4600
1535	2	51	2223	60457	1	7500
1536	1	51	2557	80750	1	5990
1537	2	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [14]: y=data1['price']  
x=data1.drop('price',axis=1)
```

In [15]:

y

```
Out[15]: 0      8900
          1      8800
          2      4200
          3      6000
          4      5700
          ...
        1533    5200
        1534    4600
        1535    7500
        1536    5990
        1537    7900
```

Name: price, Length: 1538, dtype: int64

In [16]:

x

```
Out[16]:
```

	model	engine_power	age_in_days	km	previous_owners
0	1	51	882	25000	1
1	2	51	1186	32500	1
2	3	74	4658	142228	1
3	1	51	2739	160000	1
4	2	73	3074	106880	1
...
1533	3	51	3712	115280	1
1534	1	74	3835	112000	1
1535	2	51	2223	60457	1
1536	1	51	2557	80750	1
1537	2	51	1766	54276	1

1538 rows × 5 columns

In [17]: `#!pip3 install scikit-learn`

```
In [18]: #data1=pd.get_dummies(data1)
```

```
In [19]: #data1
```

```
In [ ]:
```

```
In [20]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [21]: x_test.head(5)
```

```
Out[21]:
```

	model	engine_power	age_in_days	km	previous_owners
481	2	51	3197	120000	2
76	2	62	2101	103000	1
1502	1	51	670	32473	1
669	1	51	913	29000	1
1409	1	51	762	18800	1

```
In [22]: x_train.shape
```

```
Out[22]: (1030, 5)
```

```
In [23]: y_train.shape
```

```
Out[23]: (1030,)
```

```
In [24]: x_train.head()
```

```
Out[24]:
```

	model	engine_power	age_in_days	km	previous_owners
527	1	51	425	13111	1
129	1	51	1127	21400	1
602	2	51	2039	57039	1
331	1	51	1155	40700	1
323	1	51	425	16783	1

```
In [25]: y_train.head()
```

```
Out[25]: 527    9990
129     9500
602     7590
331     8750
323     9100
Name: price, dtype: int64
```

```
In [26]: x_test.head()
```

```
Out[26]:
```

	model	engine_power	age_in_days	km	previous_owners
481	2	51	3197	120000	2
76	2	62	2101	103000	1
1502	1	51	670	32473	1
669	1	51	913	29000	1
1409	1	51	762	18800	1

```
In [27]: y_test.head()
```

```
Out[27]: 481      7900  
        76      7900  
        1502    9400  
        669    8500  
        1409    9700  
        Name: price, dtype: int64
```

```
In [29]: from sklearn.linear_model import LinearRegression  
        reg=LinearRegression()  
        reg.fit(x_train,y_train)
```

```
Out[29]: LinearRegression()  
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.
```

```
In [ ]: #!/pip3 install scikit-learn
```

```
In [31]: ypred=reg.predict(x_test)
```

```
In [38]: ypred
```

```
Out[38]: array([ 5994.51703157,  7263.58726658,  9841.90754881,  9699.31627673,  
                10014.19892635,  9630.58715835,  9649.4499026 , 10092.9819664 ,  
                9879.19498711,  9329.19347948, 10407.2964056 ,  7716.91706011,  
                7682.89152522,  6673.95810983,  9639.42618839, 10346.53679153,  
                9366.53363673,  7707.90063494,  4727.33552438, 10428.17092937,  
                10359.87663878, 10364.84674179,  7680.16157493,  9927.58506055,  
                7127.7284177 ,  9097.51161986,  4929.31229715,  6940.60225317,  
                7794.35120591,  9600.43942019,  7319.85877519,  5224.05298205,  
                5559.52039134,  5201.35403287,  8960.11762682,  5659.72968338,  
                9915.79926869,  8255.93615893,  6270.40332834,  8556.73835062,  
                9749.72882426,  6873.76758364,  8951.72659758, 10301.95669828,  
                8674.89268564, 10301.93257222,  9165.73586068,  8846.92420399,  
                7044.68964545,  9052.4031418 ,  9390.75738772, 10267.3912561 ,  
                10046.90924744,  6855.71260655,  9761.93338967,  9450.05744337,  
                9274.98388541, 10416.00474283,  9771.10646661,  7302.96566423,  
                10082.61483093,  6996.96553454,  9829.40534825,  7134.21944391,  
                6407.26222178,  9971.82132188,  9757.01618446,  8614.84049875,  
                8437.92452169,  6489.24658616,  7752.65456507,  6626.60510856,  
                8329.88998217, 10412.00324329,  7342.77348105,  8543.63624413,  
                8706.44742777, 10010.42502651,  7256.06786062,  8522.14000051])
```

```
In [34]: from sklearn.metrics import r2_score  
         r2_score(y_test,ypred)
```

```
Out[34]: 0.8383895235218546
```

```
In [35]: from sklearn.metrics import mean_squared_error as ns  
         o=ns(y_test,ypred)  
         o
```

```
Out[35]: 593504.2888137395
```

```
In [36]: import math  
         math.sqrt(o)
```

```
Out[36]: 770.3922954013361
```


In [39]: ypred

```
Out[39]: array([ 5994.51703157,  7263.58726658,  9841.90754881,  9699.31627673,
 10014.19892635,  9630.58715835,  9649.4499026 , 10092.9819664 ,
  9879.19498711,  9329.19347948, 10407.2964056 ,  7716.91706011,
  7682.89152522,  6673.95810983,  9639.42618839, 10346.53679153,
  9366.53363673,  7707.90063494,  4727.33552438, 10428.17092937,
 10359.87663878, 10364.84674179,  7680.16157493,  9927.58506055,
  7127.7284177 ,  9097.51161986,  4929.31229715,  6940.60225317,
  7794.35120591,  9600.43942019,  7319.85877519,  5224.05298205,
  5559.52039134,  5201.35403287,  8960.11762682,  5659.72968338,
  9915.79926869,  8255.93615893,  6270.40332834,  8556.73835062,
  9749.72882426,  6873.76758364,  8951.72659758, 10301.95669828,
  8674.89268564, 10301.93257222,  9165.73586068,  8846.92420399,
  7044.68964545,  9052.4031418 ,  9390.75738772, 10267.3912561 ,
 10046.90924744,  6855.71260655,  9761.93338967,  9450.05744337,
  9274.98388541, 10416.00474283,  9771.10646661,  7302.96566423,
 10082.61483093,  6996.96553454,  9829.40534825,  7134.21944391,
  6407.26222178,  9971.82132188,  9757.01618446,  8614.84049875,
  8437.92452169,  6489.24658616,  7752.65456507,  6626.60510856,
  8329.88998217, 10412.00324329,  7342.77348105,  8543.63624413,
  6706.44742777, 10010.42502651,  7256.06706062,  8522.1400051 ])
```

```
In [42]: Results=pd.DataFrame(columns=['price', 'predicted'])
Results['price']=y_test
Results['predicted']=ypred
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(15)
```

```
Out[42]:
```

	index	price	predicted	ID
0	481	7900	5994.517032	0
1	76	7900	7263.587267	1
2	1502	9400	9841.907549	2
3	669	8500	9699.316277	3
4	1409	9700	10014.198926	4
5	1414	9900	9630.587158	5
6	1089	9900	9649.449903	6
7	1507	9950	10092.981966	7
8	970	10700	9879.194987	8
9	1198	8999	9329.193479	9
10	1088	9890	10407.296406	10
11	576	7990	7716.917060	11
12	965	7380	7682.891525	12
13	1488	6800	6673.958110	13
14	1432	8900	9639.426188	14

```
In [43]: Results['price_diff']=Results.apply(lambda row: row.price - row.predicted,axis=1)
```

In [44]: Results

Out[44]:

	index	price	predicted	ID	price_diff
0	481	7900	5994.517032	0	1905.482968
1	76	7900	7263.587267	1	636.412733
2	1502	9400	9841.907549	2	-441.907549
3	669	8500	9699.316277	3	-1199.316277
4	1409	9700	10014.198926	4	-314.198926
...
503	291	10900	10007.364639	503	892.635361
504	596	5699	6390.174715	504	-691.174715
505	1489	9500	10079.478928	505	-579.478928
506	1436	6990	8363.337585	506	-1373.337585
507	575	10900	10344.486077	507	555.513923

508 rows × 5 columns

In []: